

# Second-Generation Sequence Data Analysis with R and Bioconductor

Sean Davis

National Cancer Institute,  
National Institutes of Health  
Bethesda, MD, USA

`sdavis2@mail.nih.gov`

September 16, 2009

## Abstract

This document has two main goals: (i) To document some common and simple use cases for sequence data analysis using R and Bioconductor and (ii) To provide a playground for experimental functions for publicly (and locally) available data.

## 1 Introduction

Second-generation sequencing technologies bring huge data volumes, complex experimental design, algorithmic and data visualization challenges, and data integration nightmares. While much of the technology is proprietary, the software to deal with the data has largely been left to the open-source community. The Bioconductor project is one such open-source community that is working on problems related to second-generation sequencing. The idea is to document workflows for as many use cases as possible, not with the goal of becoming the best tool for any particular workflow, but to provide a set of tools that are useful for sequence data analysis. These tools are being developed collaboratively within the Bioconductor community.

## 2 Use Cases

While there are a huge number of applications for second-generation sequencing, there are a few use cases that demonstrate the current functionality for dealing with sequence data from within Bioconductor. The use cases will necessarily be somewhat abridged and illustrative and are not meant to be “full analyses” by any means. In fact, the power and flexibility of using Bioconductor instead of a “canned” software can really only be appreciated by extending these analyses somewhat beyond what is presented here.

Sequence analysis using R and Bioconductor relies on several “packages” that provide extended functionality beyond core R. We start by loading these packages into R:

```
> suppressMessages(library(ShortRead))  
> suppressMessages(library(Rpressa))
```

## 2.1 Targeted Sequencing

Sequencing whole genomes is still generally prohibitive in terms of time and money for most labs. Therefore, various molecular biology methods have been developed to enrich the regions of the genome of most biological interest. In the use case presented here, capture probes of 120 base pairs were used for hybridization of exons in genes of interest. Approximately 27,000 120-mer probes were used for the experiment here. The targeted genes were chosen because they have a higher likelihood of being mutated in cancer or are in pathways that are of interest in cancer.

The first step is to simply load a lane of data:

```
> data(targeted)
> targeted


class: AlignedRead
length: 6349344 reads; width: 40 cycles
chromosome: QC QC ... QC QC
position: NA NA ... NA NA
strand: NA NA ... NA NA
alignQuality: NumericQuality
alignData varLabels: run lane ... filtering contig
```

The R object `targeted` contains data on 6349344 reads including those reads that were poor quality and those that did not align to the human genome. Getting the reads that align to the human genome is fairly easy.

```
> aln2 <- targeted[!is.na(position(targeted))]
```

The data in `aln2` are now those reads that align to the human genome; there are 4995174 such reads (78.67 % of the reads). There are a number of quality control functions and accessors that could be applied to these data. An interesting one is to look at the proportion of bases at each cycle of the read.

```
> library(lattice)
> abc <- (alphabetByCycle(sread(aln2))/length(aln2))[1:4, ]
> colnames(abc) <- 1:40
> abc <- t(abc)
> matplot(abc, type = "b", xlab = "Cycle", ylab = "Proportion of Reads",
+         pch = c("A", "C", "G", "T"))
```



It is interesting to look at sequencing coverage for the capture regions as a quality control measure. Loading the description of the capture regions from a bed-format file is quite straightforward using another Bioconductor package, *rtracklayer*.

```
> suppressMessages(library(rtracklayer))
> bedfile <- system.file("extdata/agilent27k.lot1.bed", package = "Rpressa")
> rl.capture <- ranges(import(bedfile))
> rl.capture
```


```
CompressedIRangesList: 25 elements
names(25): chr1 chr10 chr11 chr12 chr13 ... chr8 chr9 chrX chrX_random chrY
```

Calculating the number of times each base in the genome has been sequenced is also easily accomplished.

```
> cvg <- coverage(aln2)
> names(cvg) <- sub(".fa", "", names(cvg))
> cvg <- cvg[names(cvg) %in% names(rl.capture)]
> vcvg <- Views(cvg, rl.capture)
```

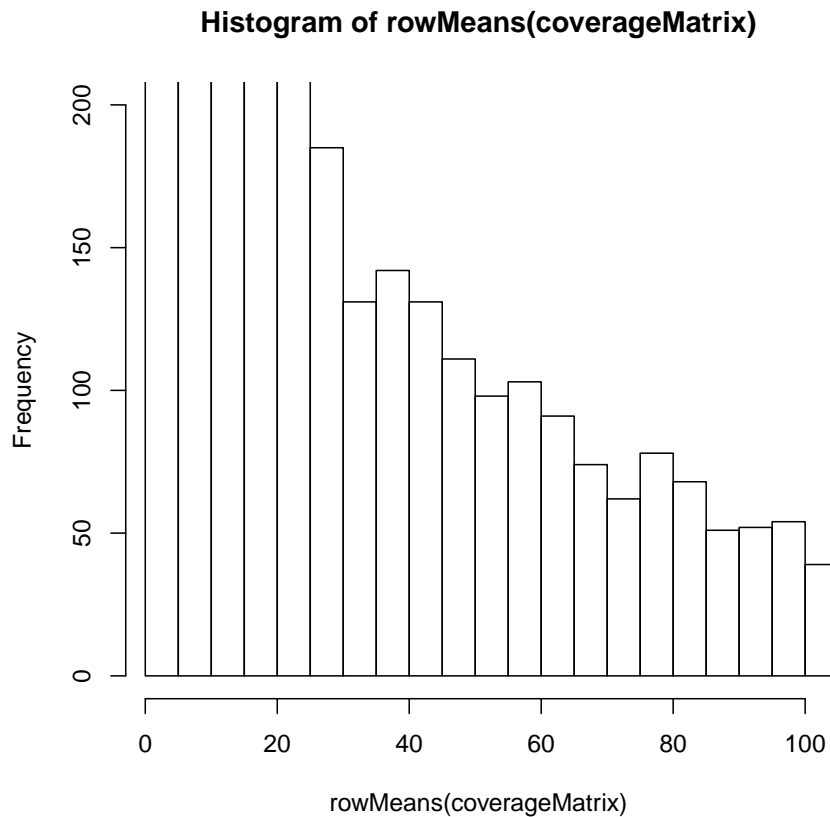
Now, *vcvg* is a view of the genomic coverage that contains the regions targeted by the targeting probes. How well are the targeted regions covered?

```
> coverageMatrix <- t(do.call(cbind, as.list(viewApply(vcvg, as.vector))))  
> plot(colMeans(coverageMatrix), type = "l", col = "red", xlab = "Base Pairs",  
+       ylab = "Coverage")
```




And how many regions have a mean coverage above a given threshold?

```
> hist(rowMeans(coverageMatrix), breaks = "scott", xlim = c(1,  
+       100), ylim = c(0, 200))
```



In short, 7866 regions have mean coverage > 1 while 3053 have coverage > 10. Bioconductor also has data packages that contain the entire human genome sequence, but in a compact and random-accessible form for memory efficiency. With this information in hand, it might be interesting to look at the effect of the GC content of the capture oligos on genomic coverage.

```
> library(BSgenome)
> library(BSgenome.Hsapiens.UCSC.hg18)
> regionDNA <- DNASTringSet(getSeq(Hsapiens, rl.capture))
> x <- alphabetFrequency(regionDNA, as.prob = TRUE)
> avgCvg <- rowMeans(coverageMatrix)
> x <- x[avgCvg > 0, 2:3]
> plot(rowSums(x), avgCvg[avgCvg > 0], log = "xy", xlab = "GC content",
+       ylab = "Coverage", pch = ".")
```



## 2.2 ChIP-Seq

This section borrows heavily from the vignette for the *chipseq*.

### Example data

The *cstest* data set is included in the *chipseq* package to help demonstrate its capabilities. The dataset contains data for three chromosomes from Solexa lanes, one from a CTCF mouse ChIP-Seq, and one from a GFP mouse ChIP-Seq. The raw reads were aligned to the reference genome (mouse in this case) using an external program (MAQ), and the results read in using the `readReads` function, which in turn uses the `readAligned` function in the *ShortRead*. This step removed all duplicate reads and applied a quality score cutoff. The remaining data were reduced to a set of alignment start positions (including orientation).

```
> suppressMessages(library(chipseq))
> data(cstest)
> cstest
```

```
GenomeDataList: 2 elements
names(2): ctcf gfp
```

## Extending Reads

The sequencer generally reads only the first  $n$  (where  $n$  is typically on the order of 36bp for ChIP-seq applications), but the typical insert size is on the order of 150-250bp. The TFBS of interest is somewhere in that fragment (ideally), but not necessarily at the beginning. Therefore, it is useful to extend the reads to the full length of the average fragment size. To make the chromosomes the correct size, the lengths of the chromosomes need to be used.

```
> library(BSgenome.Mmusculus.UCSC.mm9)
> mouse.chromlens <- seqlengths(Mmusculus)
> head(mouse.chromlens)

      chr1      chr2      chr3      chr4      chr5      chr6
197195432 181748087 159599783 155630120 152537259 149517037
```

We extend all reads to be 200 bases long. This is done using the `extendReads()` function, which can work on data from one chromosome in one lane.

```
> ext <- extendReads(cstest$ctcf$chr10, seqLen = 200)
> head(ext)
```

```
IRanges instance:
      start      end width
[1] 3012936 3013135    200
[2] 3012941 3013140    200
[3] 3012944 3013143    200
[4] 3012955 3013154    200
[5] 3012963 3013162    200
[6] 3012969 3013168    200
```

As with the targeted sequencing example, computing coverage can be a useful way of looking at the data. To keep things simple, the analysis is first restricted to chromosome 10.

```
> cov <- coverage(ext, width = mouse.chromlens["chr10"])
> cov

'integer' Rle instance of length 129993255 with 288928 runs
Lengths:  3012799 97 2 37 5 3 11 8 6 9 ...
Values :  0 1 2 3 5 6 7 8 9 10 ...
```

The regions of interest are contiguous segments of non-zero coverage, which the *chipseq* package refers to as *islands*.

```
> islands <- slice(cov, lower = 1)
> islands
```

Views on a 129993255-length Rle subject

```
views:
      start      end width
[1] 3012800 3013270   471 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
[2] 3018464 3018663   200 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
```

```

[3] 3020766 3020965 200 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
[4] 3023019 3023218 200 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
[5] 3023240 3023439 200 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
[6] 3032536 3032735 200 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
[7] 3038377 3038576 200 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
[8] 3040312 3040554 243 [1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 ...]
[9] 3041098 3041297 200 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
...
[87957] 129973175 129973447 273 [1 1 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 ...]
[87958] 129974813 129975012 200 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
[87959] 129975575 129975774 200 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
[87960] 129978669 129978868 200 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
[87961] 129979209 129979571 363 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
[87962] 129980253 129980452 200 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
[87963] 129981957 129982156 200 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
[87964] 129982330 129982529 200 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]
[87965] 129987020 129987219 200 [1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 ...]

```

For each island, we can compute the number of reads in the island, and the maximum coverage depth within that island.

```

> viewSums(head(islands))

[1] 2400 200 200 200 200 200

> viewMaxs(head(islands))

[1] 11 1 1 1 1 1

> nread.tab <- table(viewSums(islands)/200)
> depth.tab <- table(viewMaxs(islands))
> head(nread.tab, 10)

 1    2    3    4    5    6    7    8    9   10
68111 13350 3022  925  415  247  191  122  132  100

> head(depth.tab, 10)

 1    2    3    4    5    6    7    8    9   10
68159 14745 2388  547  256  180  150  129  120  102

```

It is also possible to process all the data in all lanes simultaneously.

```

> islandReadSummary <- function(x) {
+   g <- extendReads(x, seqLen = 200)
+   s <- slice(coverage(g), lower = 1)
+   tab <- table(viewSums(s)/200)
+   ans <- data.frame(nread = as.numeric(names(tab)), count = as.numeric(tab))
+   ans
+ }

```



```

> nread.islands <- gdapply(cstest, islandReadSummary)
> nread.islands <- as(nread.islands, "data.frame")
> head(nread.islands)

```

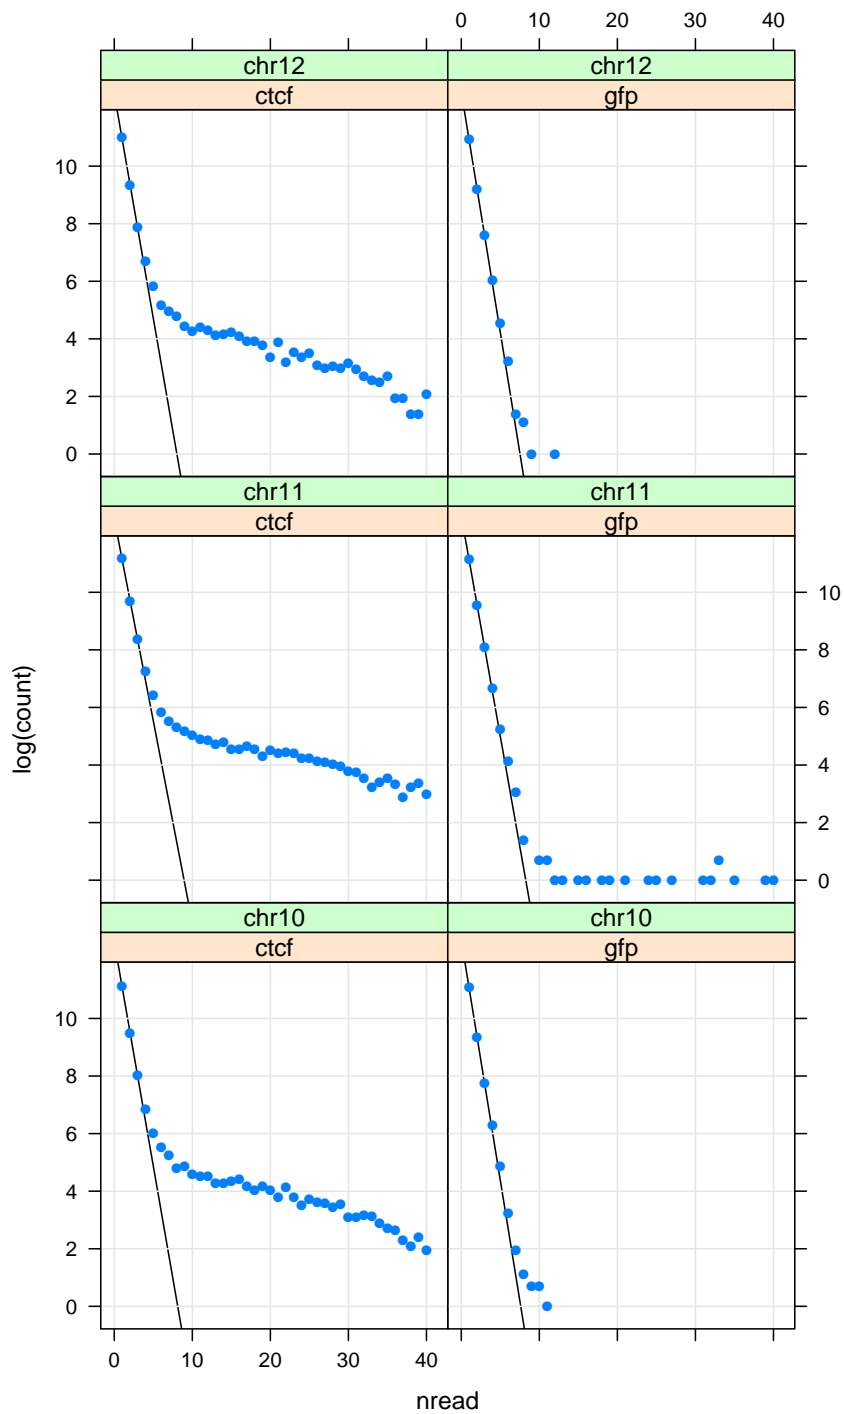
	nread	count	chromosome	sample
1	1	68111	chr10	ctcf
2	2	13350	chr10	ctcf
3	3	3022	chr10	ctcf
4	4	925	chr10	ctcf
5	5	415	chr10	ctcf
6	6	247	chr10	ctcf

A simple plot of the  $\log(\text{count})$  versus the number of reads in each island is useful. If the reads were randomly distributed across the genome, the relationship should be linear. In the GFP lane, this is close to true for most of the data. However, for CTCF, there is obvious deviation from linear. Points to the right of the line in each plot are, then, “significant” in some sense and a threshold of 8 reads in an island looks like a good pick for finding islands of significance.

```

> xyplot(log(count) ~ nread | sample + chromosome, nread.islands,
+       subset = (nread <= 40), pch = 16, type = c("p", "g"), panel = function(x,
+       y, ...) {
+       panel.lmline(x[1:3], y[1:3], col = "black")
+       panel.xyplot(x, y, ...)
+     })

```



Finding peaks is a fairly simple procedure. Note that the threshold defined by the linear extrapolation in the plot is used.

```

> peaks <- slice(cov, lower = 8)
> peaks

Views on a 129993255-length Rle subject


views:
      start      end width
[1] 3012955 3013135 181 [ 8 8 8 8 8 8 8 8 8 9 9 9 9 9 ...]
[2] 3234799 3234895 97 [ 8 8 8 8 8 8 8 8 8 8 8 8 8 8 ...]
[3] 3270012 3270297 286 [8 8 8 9 9 9 9 9 8 8 8 8 8 8 8 8 8 ...]
[4] 3277662 3277832 171 [ 8 8 8 8 8 8 8 8 8 8 8 8 8 8 ...]
[5] 3277848 3277859 12 [8 8 8 8 8 8 8 8 8 8 8 8]
[6] 3460857 3460973 117 [ 8 8 8 8 8 8 8 8 8 8 8 8 8 8 ...]
[7] 3617850 3617983 134 [ 8 8 8 8 8 8 8 9 9 9 9 10 10 ...]
[8] 3651712 3651992 281 [ 8 8 9 9 9 10 10 10 10 10 10 10 10 ...]
[9] 4310402 4310712 311 [ 8 8 8 9 9 9 9 9 9 9 9 9 10 10 ...]
...
[1747] 128986519 128986595 77 [8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 ...]
[1748] 128986604 128986610 7 [8 8 8 8 8 8]
[1749] 128986638 128986673 36 [8 8 8 8 8 8 9 9 9 9 8 8 8 9 9 9 8 8 8 ...]
[1750] 129058889 129058980 92 [8 8 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9 ...]
[1751] 129530031 129530201 171 [ 8 8 9 9 9 9 9 9 9 9 9 9 9 9 9 ...]
[1752] 129533303 129533381 79 [8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 ...]
[1753] 129665351 129665586 236 [ 8 9 9 9 9 9 9 9 9 9 10 10 10 10 ...]
[1754] 129666784 129666947 164 [8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8 ...]
[1755] 129750671 129750849 179 [ 8 8 8 8 8 8 9 9 9 9 9 9 9 9 ...]

> peak.depths <- viewMaxs(peaks)
> cov.pos <- coverage(extendReads(cstest$ctcf$chr10, strand = "+",
+   seqLen = 200), width = mouse.chromlens["chr10"])
> cov.neg <- coverage(extendReads(cstest$ctcf$chr10, strand = "-",
+   seqLen = 200), width = mouse.chromlens["chr10"])
> peaks.pos <- copyIRanges(peaks, cov.pos)
> peaks.neg <- copyIRanges(peaks, cov.neg)
> wpeaks <- tail(order(peak.depths), 4)
> wpeaks

[1] 971 989 1079 922

> coverageplot(peaks.pos[wpeaks[1]], peaks.neg[wpeaks[1]])

```



```
> coverageplot(peaks.pos[wpeaks[2]], peaks.neg[wpeaks[2]])
```

